

SRI KRISHNA N

AI / ML Engineer

+1 (972) 695-9606 • srikrishna3099@gmail.com

PROFESSIONAL SUMMARY

AI Engineer with 5+ years of experience developing and deploying end-to-end ML and GenAI solutions for large enterprises. Deep expertise in LLMs, RAG pipelines, LangChain, agentic AI workflows, and production-grade deep learning systems. Proven ability to deploy scalable GenAI applications using AWS Bedrock and Hugging Face models, collaborate with engineering teams to build robust AI APIs, and promote responsible AI best practices. Experienced in probabilistic modeling, NLP, and model lifecycle management across cloud-native environments (AWS, Azure, GCP) with Docker and Kubernetes.

TECHNICAL SKILLS

ML Frameworks	PyTorch, TensorFlow, Keras, Scikit-learn, Spark MLlib, Hugging Face Transformers
Deep Learning	CNNs, RNNs, LSTMs, GANs, Transformers, BERT, GPT fine-tuning, Diffusion Models
GenAI & LLM	LangChain, LlamaIndex, RAG Pipelines, Tool Use / Function Calling, Prompt Engineering, LLM Evaluation, Hugging Face, AWS Bedrock, Azure OpenAI
MLOps & DevOps	Docker, Kubernetes, MLflow, CI/CD (Jenkins, GitHub Actions), Flask, FastAPI
Cloud Platforms	AWS (SageMaker, Lambda, S3), GCP (Vertex AI, BigQuery), Azure ML Studio
Big Data	Apache Spark, Hadoop, Hive, Apache Kafka, Airflow
Languages	Python, R, SQL, Scala
Databases	PostgreSQL, MySQL, MongoDB, Cassandra, HBase, FAISS, Pinecone (Vector DBs)
Visualization	Tableau, Power BI, Matplotlib, Seaborn

WORK EXPERIENCE

Senior AI/ML Engineer

Jun 2024 – Present

AT&T | Cleveland, Ohio, USA

- Designed and deployed end-to-end RAG pipelines using LangChain and AWS Bedrock, enabling enterprise scale semantic search and document Q&A over proprietary datasets significantly reducing manual information retrieval time across business units.
- Built multi-step agentic AI workflows using LangChain tool-use and function calling, orchestrating LLM-driven automation pipelines that integrated with internal APIs and enterprise data sources for intelligent decision support.
- Deployed scalable GenAI applications backed by Hugging Face open-source models and AWS Bedrock foundation models, building production-grade REST APIs using FastAPI and Flask containerized with Docker and Kubernetes.
- Engineered production ML/AI solutions for predictive analytics and customer segmentation using PyTorch, TensorFlow, and Azure ML optimizing hyperparameters via Bayesian search and cutting model inference latency.
- Designed and maintained full ML model lifecycle: automated CI/CD pipelines (GitHub Actions, Jenkins), containerized serving (Docker, Kubernetes), model monitoring, data drift detection, and automated retraining workflows.
- Applied probabilistic modeling, A/B testing, and statistical validation to evaluate LLM and ML model performance, maintaining rigorous benchmarks across foundation models and fine-tuned systems.
- Collaborated with cross-functional software engineers and product managers in Agile sprints to scope and deliver AI-powered applications, reducing customer churn by 35% and time-to-deployment for new AI features.
- Implemented SHAP and LIME explainability frameworks and established responsible AI deployment standards, increasing stakeholder trust and ensuring compliance with enterprise AI governance requirements.

Senior Data Scientist / ML Engineer

Nov 2023 – May 2024

Procter & Gamble | Cincinnati, Ohio, USA

- ▶ Applied NLP and deep learning (Transformers, TF-IDF embeddings, Hugging Face models) to extract actionable business insights from unstructured text; utilized PySpark MLlib for scalable feature engineering and model training.
- ▶ Built real-time data streaming pipelines with Apache Kafka and orchestrated end-to-end ML workflows using Apache Airflow, reducing pipeline failures by 40% and improving data freshness.
- ▶ Leveraged Azure Databricks and Azure ML to develop and optimize predictive models for supply chain and demand forecasting, reducing forecast error by 25% by automating hyperparameter tuning.
- ▶ Designed ETL pipelines with Hive and Pig; implemented Random Forest, Decision Tree, and SVM models for package classification with high accuracy, evaluated using ROC/AUC metrics.
- ▶ Conducted model performance benchmarking and version control using MLflow, enabling consistent experiment tracking and rapid iteration across the ML team.
- ▶ Developed multi-class text classification pipelines using TF-IDF and transformer-based embeddings (Hugging Face) for automated product quality tagging, prototyping AI-powered applications that reduced manual review effort significantly.

Machine Learning Engineer

Aug 2022 – Jul 2023

Fortis Healthcare | Hyderabad, India

- ▶ Led end-to-end ML lifecycle for clinical predictive models, data collection, feature engineering, training, validation, and cloud-native deployment using Docker, Kubernetes, and AWS Lambda for scalable real-time inference.
- ▶ Developed automated model retraining pipelines to continuously update clinical prediction models with new patient data, ensuring sustained performance over time.
- ▶ Designed scalable ETL pipelines using Apache Spark and Informatica PowerCenter, processing 1M+ patient EHR records for analytics and reporting in compliance with healthcare data standards.
- ▶ Built and validated clinical dashboards in Tableau, implementing incremental refresh and automated data pipelines to enable real-time decision-making for medical professionals.
- ▶ Implemented Change Data Capture (CDC) to preserve historical data integrity in SCD tables and designed optimized database schemas with indexing strategies for faster clinical data retrieval.
- ▶ Applied hypothesis testing and power analysis in Python to validate clinical model outputs, reducing manual review time by 40% and ensuring statistical reliability before deploying models to production.

Data Scientist

Mar 2021 – Jul 2022

Standard Chartered Bank | Hyderabad, India

- ▶ Developed and deployed fraud detection models using XGBoost, Random Forest, and Logistic Regression, improving detection accuracy and leading to a 20% increase in customer acquisition.
- ▶ Applied NLP and ML to extract insights from unstructured banking data (customer feedback, transaction records) using Apache Spark and Hive for distributed text processing.
- ▶ Automated ML model training and hyperparameter tuning with Azure AutoML, reducing development time and improving forecasting accuracy for financial risk models.
- ▶ Engineered features from raw transactional data using Pandas and NumPy, improving signal quality for downstream risk scoring models and reducing false positive fraud alerts.
- ▶ Presented model findings and risk insights to senior banking stakeholders through clear data storytelling and Tableau dashboards, enabling faster executive decision-making on credit policy adjustments.

EDUCATION

Master's in Information Technology

GPA: 3.8 / 4.0

University of Cincinnati, Ohio, USA

PUBLICATION

Machine Learning and Statistical Techniques for Outlier Detection in Smart Home Energy Consumption — 2024 IEEE Open Conference of Electrical, Electronic and Information Sciences (eStream). DOI: 10.1109/eStream62072.2024.10546900